

# Lab2

September 30, 2021

```
[ ]: import numpy as np
import pandas as pd
import sklearn
from sklearn.metrics import confusion_matrix, accuracy_score, make_scorer
from sklearn.model_selection import cross_validate, cross_val_score, □
    ↳cross_val_predict, GridSearchCV, RandomizedSearchCV
from sklearn.neighbors import KNeighborsClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.naive_bayes import GaussianNB
from sklearn.preprocessing import StandardScaler
from sklearn.svm import SVC
import matplotlib
import matplotlib.pyplot as plt
import warnings

warnings.simplefilter(action='ignore', category=pd.errors.PerformanceWarning)
```

## 1 Assignment 1.1

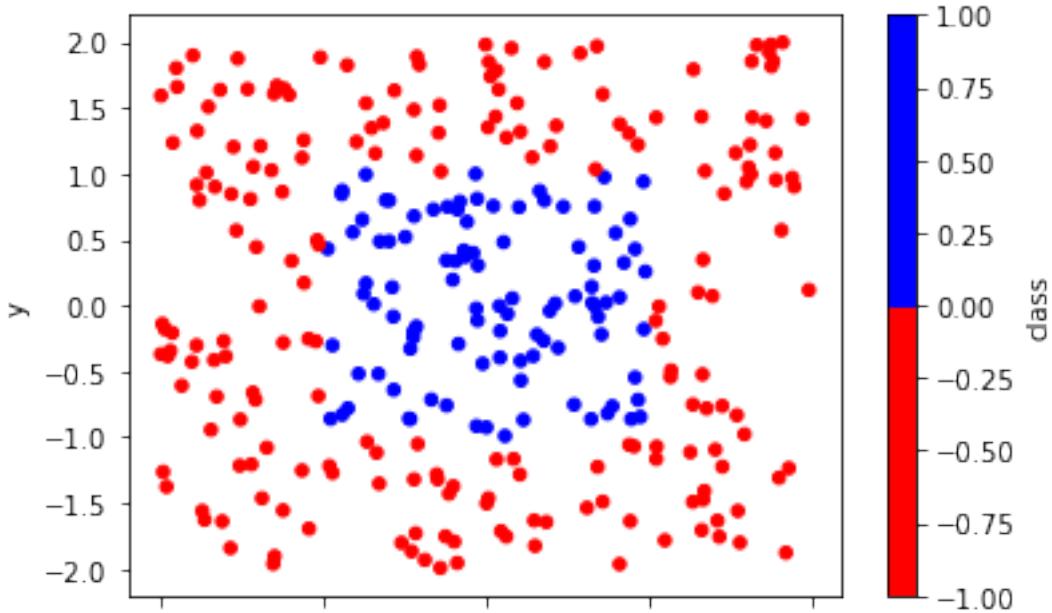
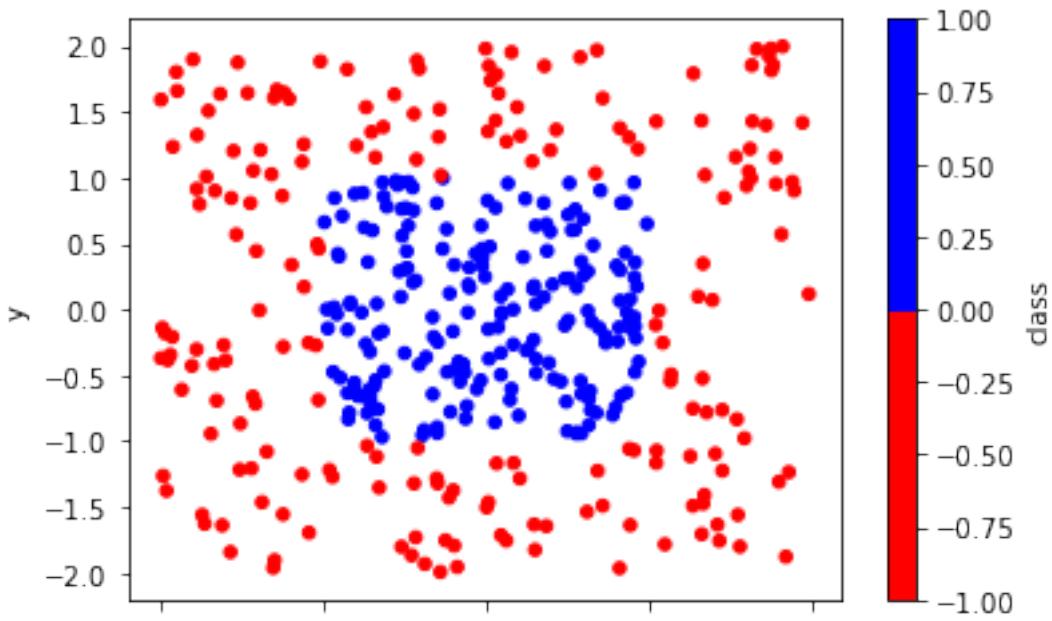
```
[ ]: exp1a = pd.read_csv('exp1a.csv')
exp1a.name = 'exp1a'
exp1b = pd.read_csv('exp1b.csv')
exp1b.name = 'exp1b'
exp1c = pd.read_csv('exp1c.csv')
exp1c.name = 'exp1c'
```

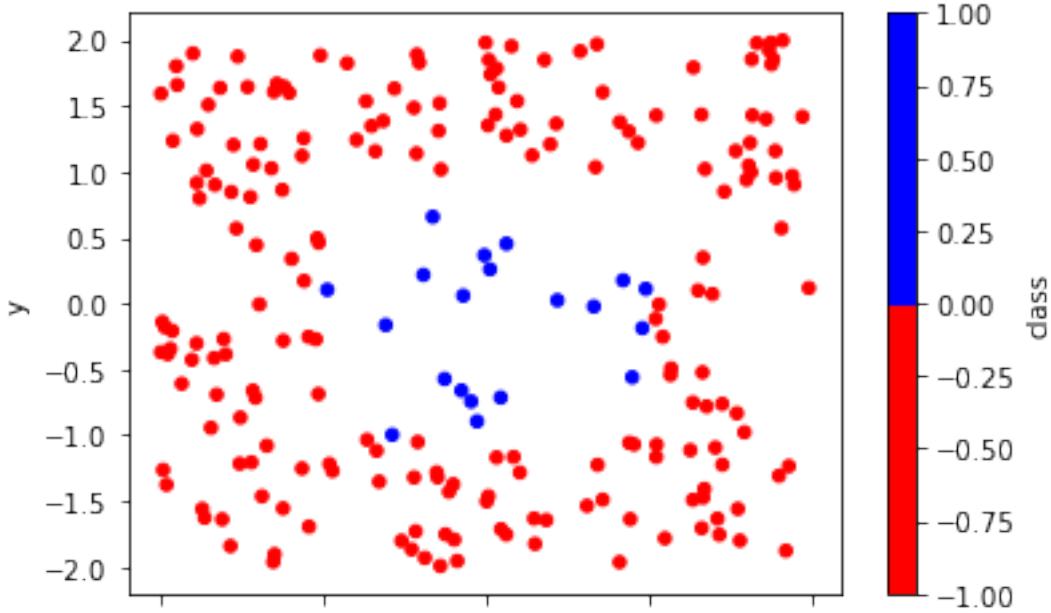
### 1.1 (a)

Visualize the data sets from Experiment 1

```
[ ]: cmap = matplotlib.colors.ListedColormap(['red', 'blue'])
exp1a.plot.scatter('x', 'y', c='class', cmap=cmap)
exp1b.plot.scatter('x', 'y', c='class', cmap=cmap)
exp1c.plot.scatter('x', 'y', c='class', cmap=cmap)
```

```
[ ]: <matplotlib.axes._subplots.AxesSubplot at 0x7fe19a995f10>
```





## 1.2 (b) and (c)

For each of the data sets, how does the 10-fold cross-validation accuracy rates and confusion matrices vary as k increases? Explain this trend.

```
[ ]: idx = [
    np.array([1, 1, 5, 5, 11, 11, 21, 21]),
    np.array(['positive', 'negative'] * 4)
]

cols = [
    np.concatenate([[i] * 3 for i in ['exp1a', 'exp1b', 'exp1c']]),
    ['confusion matrix', 'confusion matrix', 'accuracy'] * 3,
    ['positive', 'negative', ''] * 3
]
results_b_c = pd.DataFrame(columns=cols, index=idx)

results_b_c.index.names = ['k', 'predicted']
for exp in [exp1a, exp1b, exp1c]:
    for k in [1, 5, 11, 21]:
        knn = KNeighborsClassifier(n_neighbors=k)
        y_predict = cross_val_predict(knn, exp[['x', 'y']].values, □
        →exp['class'], cv=10)
        results_b_c.loc[k, (exp.name, 'confusion matrix')] = □
        →confusion_matrix(exp['class'], y_predict)
```

```

    results_b_c.loc[k,(exp.name, 'accuracy')] = u
    ↪accuracy_score(exp['class'], y_predict)
results_b_c

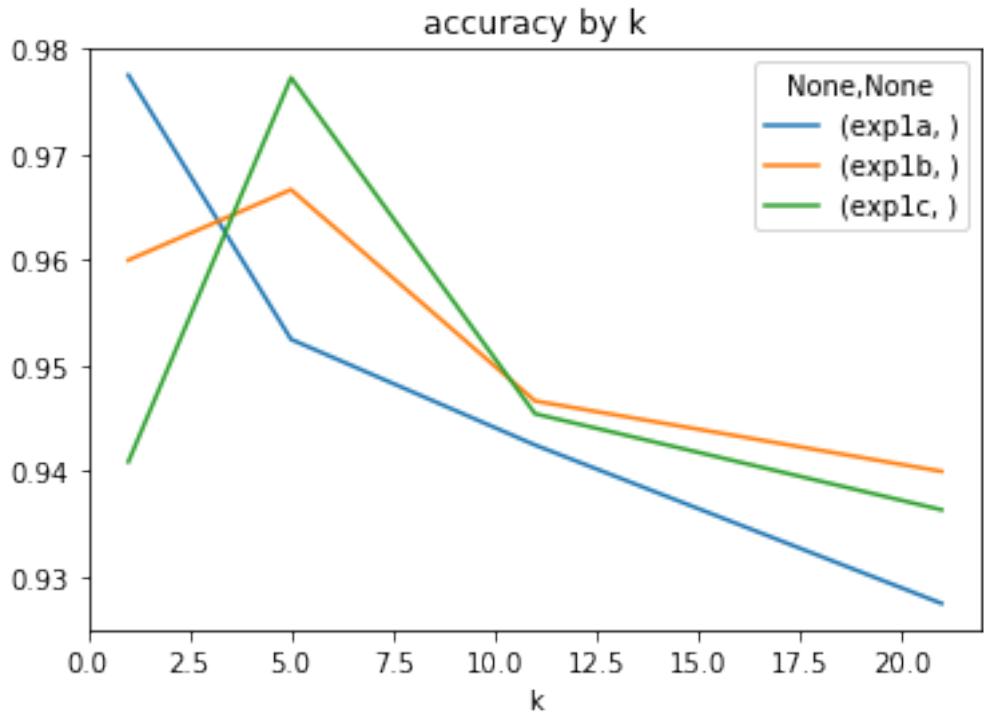
```

```
[ ]: exp1a
      confusion matrix           exp1b
      accuracy confusion matrix   \
      positive negative          positive negative
k predicted
1  positive       194        6  0.9775       193        7
   negative        3        197  0.9775        5        95
5  positive       183        17  0.9525       194        6
   negative        2        198  0.9525        4        96
11 positive      177        23  0.9425       190       10
   negative        0        200  0.9425        6        94
21 positive      171        29  0.9275       190       10
   negative        0        200  0.9275        8        92

      exp1c
      accuracy confusion matrix           accuracy
      positive negative
k predicted
1  positive      0.96       194        6  0.940909
   negative      0.96        7        13  0.940909
5  positive      0.966667     200        0  0.977273
   negative      0.966667       5        15  0.977273
11 positive      0.946667     200        0  0.945455
   negative      0.946667       12        8  0.945455
21 positive      0.94       200        0  0.936364
   negative      0.94        14        6  0.936364
```

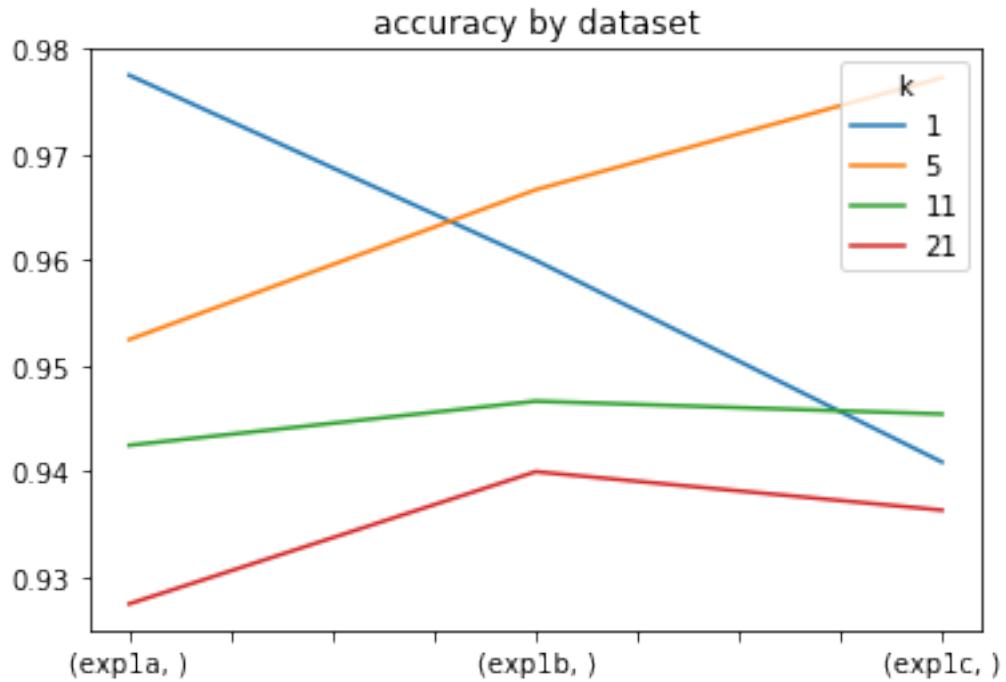
```
[ ]: (results_b_c
      .select_dtypes(object)
      .astype(float)
      .xs('accuracy', axis=1, level=1)
      .groupby('k', level=0, axis=0)
      .agg('mean')
).plot.line(title='accuracy by k')
```

```
[ ]: <matplotlib.axes._subplots.AxesSubplot at 0x7fe19adc4b80>
```



```
[ ]: (results_b_c
      .select_dtypes(object)
      .astype(float)
      .xs('accuracy', axis=1, level=1)
      .groupby('k', level=0, axis=0)
      .agg('mean')
    ).T.plot.line(title='accuracy by dataset')
```

```
[ ]: <matplotlib.axes._subplots.AxesSubplot at 0x7fe19af647c0>
```



As we can see with increasing  $k$  the accuracy drops for all data sets. However, it's important to notice, that for set  $\text{exp1b}$  and  $\text{exp1c}$  there is a spike for  $k=5$ . This can be explained by the fact that with increasing  $k$  we decrease the flexibility of the model and hence increase the bias (while decreasing the variance). According to the recorded accuracy values it can be assumed that the best tradeoff between bias and variance is at about  $k=5$ . It's also interesting to see, that the different data sets yield different results for different  $k$  values. From the chart 'accuracy by dataset' we can clearly see, the different  $k$ -values yield different results for different data sets. Whereas  $k=11$  performs only third best for  $\text{exp1a}$  and  $\text{exp1b}$  it performs second best for  $\text{exp1c}$ .

### 1.3 (d)

```
[ ]: idx = [
    np.array([1, 1, 5, 5, 11, 11, 21, 21]),
    np.array(['positive', 'negative'] * 4)
]

cols = [
    np.concatenate([[i] * 3 for i in ['exp1a', 'exp1b', 'exp1c']]),
    ['confusion matrix', 'confusion matrix', 'accuracy'] * 3,
    ['positive', 'negative', ''] * 3
]
results_d = pd.DataFrame(columns=cols, index=idx)
results_d.index.names = ['k', 'predicted']
for exp in [exp1a, exp1b, exp1c]:
    for k in [1,5,11,21]:
```

```

knn = KNeighborsClassifier(n_neighbors=k, weights='distance')
y_predict = cross_val_predict(knn, exp[['x', 'y']].values, exp['class'].values, cv=10)
results_d.loc[k,(exp.name, 'confusion matrix')] = confusion_matrix(exp['class'], y_predict)
results_d.loc[k,(exp.name, 'accuracy')] = accuracy_score(exp['class'], y_predict)
results_d

```

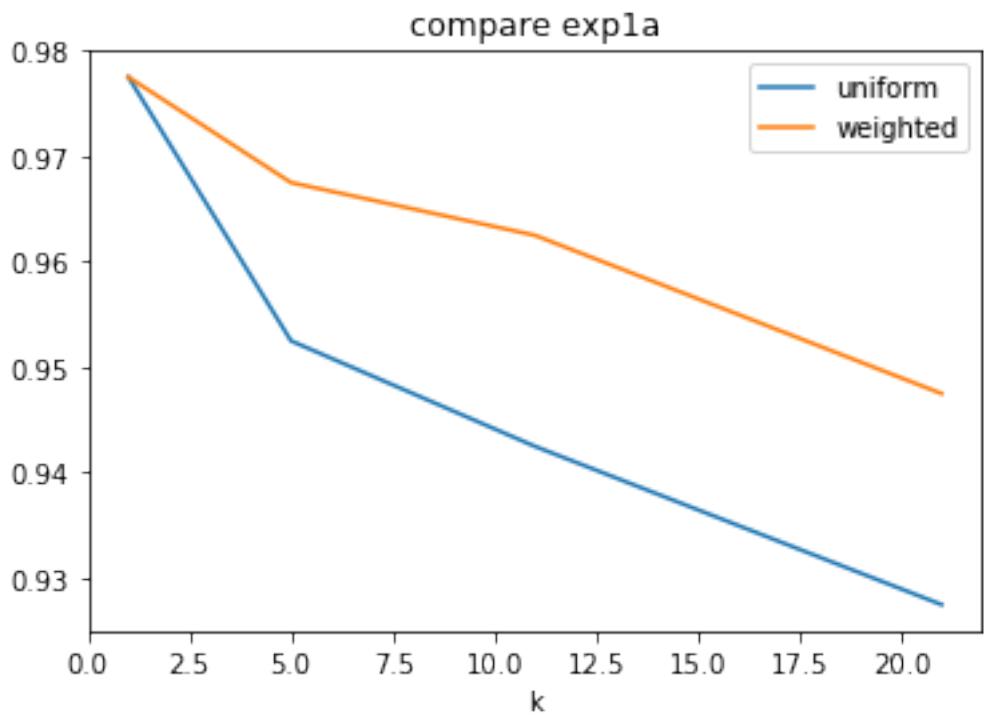
```
[ ]: exp1a
      confusion matrix           accuracy confusion matrix
      positive negative           positive negative
k predicted
1 positive       194        6  0.9775          193        7
      negative      3        197  0.9775          5        95
5 positive       189       11  0.9675          195        5
      negative      2        198  0.9675          4        96
11 positive      186       14  0.9625          194        6
      negative      1        199  0.9625          6        94
21 positive      179       21  0.9475          193        7
      negative      0        200  0.9475          4        96

      exp1b
      accuracy confusion matrix           accuracy confusion matrix
      positive negative           positive negative
k predicted
1 positive       0.96        194        6  0.940909
      negative      0.96        7        13  0.940909
5 positive       0.97        200        0  0.977273
      negative      0.97        5        15  0.977273
11 positive      0.96        200        0  0.968182
      negative      0.96        7        13  0.968182
21 positive      0.963333     200        0  0.968182
      negative      0.963333     7        13  0.968182
```

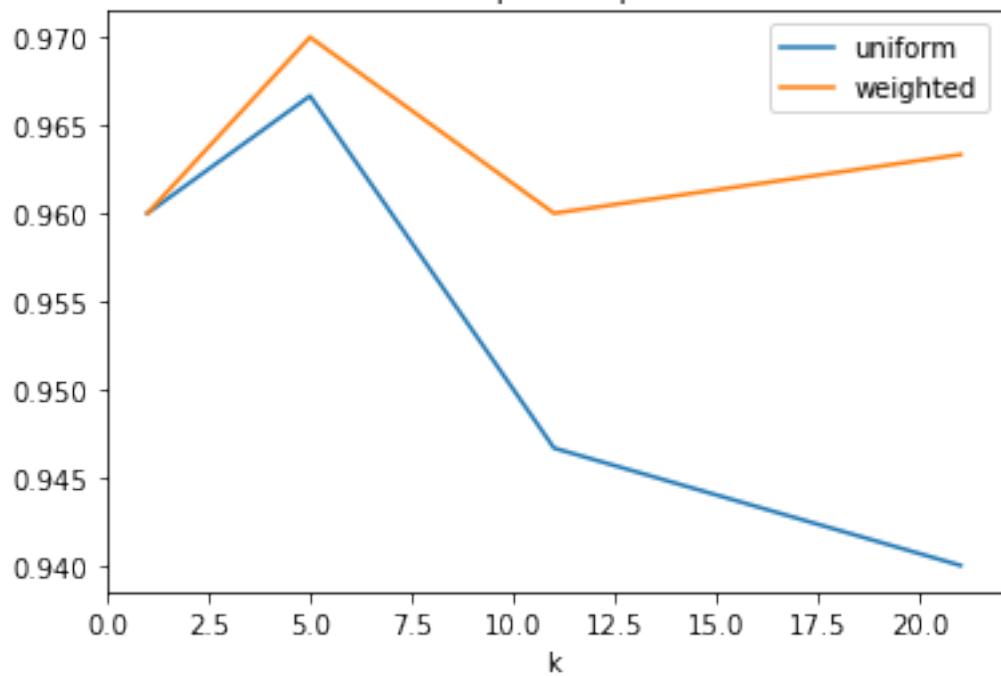
```
[ ]: compare_results = pd.concat(
    (results_b_c.select_dtypes(object).astype(float).xs('accuracy', axis=1, level=1).groupby('k', level=0, axis=0).agg('mean'),
     results_d.select_dtypes(object).astype(float).xs('accuracy', axis=1, level=1).groupby('k', level=0, axis=0).agg('mean'))
    )
    , axis=1
)
compare_results.columns = ['exp1a', 'exp1b', 'exp1c', 'exp1a_d', 'exp1b_d', 'exp1c_d']

for i in ['a', 'b', 'c']:
```

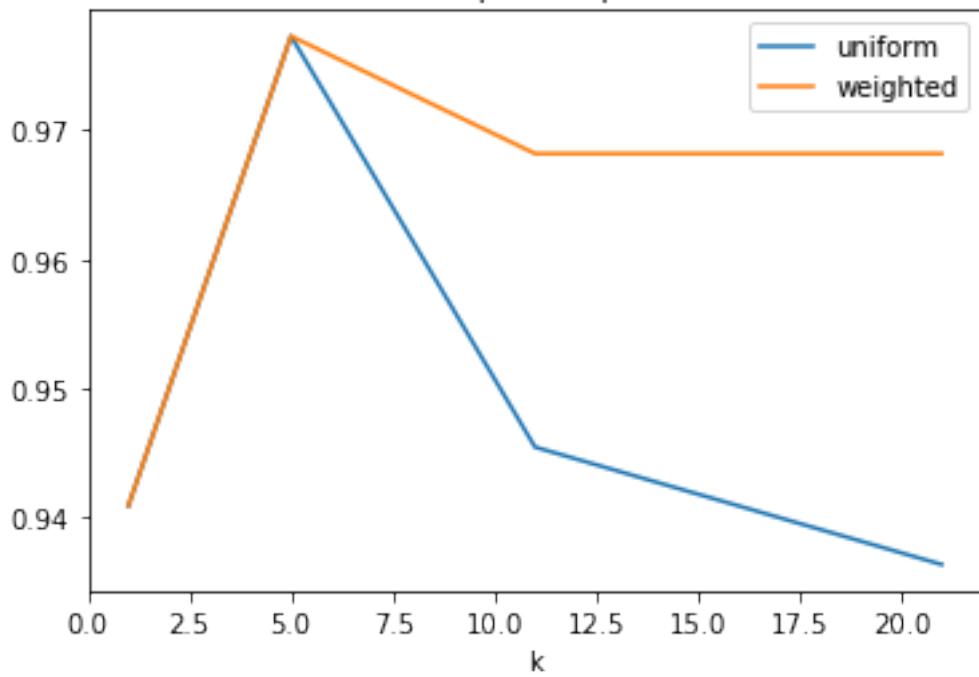
```
temp = compare_results[[f'exp1{i}', f'exp1{i}_d']]
temp.columns = ['uniform', 'weighted']
temp.plot.line(title=f'compare exp1{i}')
```



compare exp1b



compare exp1c



In the plot ‘compare results from b and d’ we can see that the classifier using the weighted kNN approach perform better (or equal for exp1c) for all values of k. For the ‘uniform’ approach (where all neighbors are weighted equally) we see a bad generalization performance for increasing k. The ‘weighted’ approach on the other hand performs better for greater k. It is assumed, that this is because close neighbors are better indicators for the class than neighbors with a greater distance. Additionally we have to consider that these data sets are noise free. So for this example it can be concluded, that weighted distances generalize better.

## 1.4 (e)

run decision trees and logistic regression on the data sets

```
[ ]: cols = [
    ['exp1a', 'exp1a', 'exp1a', 'exp1b', 'exp1b', 'exp1b', 'exp1c', 'exp1c', 'exp1c'],
    ['confusion matrix', 'confusion matrix', 'accuracy'] * 3,
    ['positive', 'negative', ''] * 3
]

idx = [
    np.concatenate(([([i]*2) for i in ['Decision Trees', 'Logistic',
    'Regression']]), axis=0),
    # np.array(['decision tree', 'decision tree', 'logistic regression',
    # 'logistic regression']),
    np.array(['positive', 'negative']) * 2
]

results_e = pd.DataFrame(columns=cols, index = idx)

results_e.index.names = ['classifier', 'predicted']
for exp in [exp1a, exp1b, exp1c]:
    des_tree = DecisionTreeClassifier(criterion='entropy')
    log_reg = LogisticRegression()
    y_predict_tree = cross_val_predict(des_tree, exp[['x', 'y']], exp['class'])
    y_predict_log_reg = cross_val_predict(log_reg, exp[['x', 'y']], exp['class'])
    results_e.loc['Decision Trees',(exp.name, 'confusion matrix')] = confusion_matrix(exp['class'], y_predict_tree)
    results_e.loc['Logistic Regression',(exp.name, 'confusion matrix')] = confusion_matrix(exp['class'], y_predict_log_reg)
    results_e.loc['Decision Trees',(exp.name, 'accuracy')] = accuracy_score(exp['class'], y_predict_tree)
    results_e.loc['Logistic Regression',(exp.name, 'accuracy')] = accuracy_score(exp['class'], y_predict_log_reg)
results_e
```

```
[ ]: exp1a
confusion matrix
accuracy \
```

		positive negative		
classifier	predicted	positive	negative	
Decision Trees	positive	200	0	0.995
	negative	2	198	0.995
Logistic Regression	positive	104	96	0.5275
	negative	93	107	0.5275

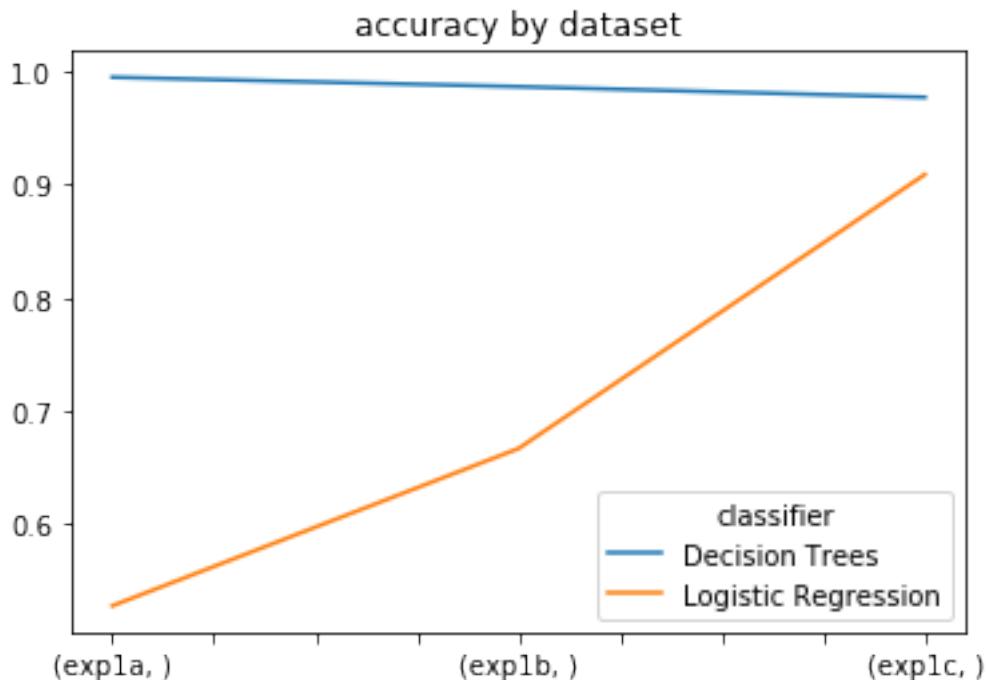
		exp1b		\
classifier	predicted	confusion matrix		accuracy
		positive	negative	
Decision Trees	positive	200	0	0.986667
	negative	4	96	0.986667
Logistic Regression	positive	200	0	0.666667
	negative	100	0	0.666667

		exp1c		
classifier	predicted	confusion matrix		accuracy
		positive	negative	
Decision Trees	positive	200	0	0.977273
	negative	5	15	0.977273
Logistic Regression	positive	200	0	0.909091
	negative	20	0	0.909091

```
[ ]: (results_e
      .select_dtypes(object)
      .astype(float)
      .xs('accuracy', axis=1, level=1)
      .groupby('classifier', level=0, axis=0)
      .agg('mean')
    ).T.plot.line(title='accuracy by dataset')
```

```
[ ]: <matplotlib.axes._subplots.AxesSubplot at 0x7fe19b307d00>
```



The data in experiment 1 does not follow a linear distribution. Hence, the logistic regression yields worse results than the decision trees. Decision trees work very well on this type of data, since the decision boundaries are parallel to the axis of the coordinate system. When looking at the chart ‘accuracy by dataset’ we can see, that for the exp1c we get way better results than for exp1a and exp1b. This is probably because the exp1c is the most imbalanced one. Whereas exp1a is split in equal proportions. Hence, we receive for exp1a an accuracy score below 0.6. This underscores the impact an unbalanced dataset can have on the classification accuracy score.

## 2 Assignment 1.2

to improve the performance of SVC/logistic regression and get fitting and prediction results in a reasonable amount of time I implemented the standard scaler. This is especially important for the exp2b data set.

```
[ ]: exp2a = pd.read_csv('exp2a.csv')
exp2a.name = 'exp2a'
exp2b = pd.read_csv('exp2b.csv')
exp2b.name = 'exp2b'
```

```
[ ]: des_tree = DecisionTreeClassifier(criterion='entropy')
knn = KNeighborsClassifier()
gau_nav_bay = GaussianNB()
log_reg = LogisticRegression()
svm_clf = SVC(kernel='linear')
```

## 2.1 (a)

which classifiers have a good, which have a bad gernalization performance?

```
[ ]: def get_evaluation(data_set, classifier_):
    classifier_names = [str(j.__class__).split('.')[1][-1][:-2] for j in classifier_]
    cols = [
        np.concatenate([(i.name]*3) for i in [data_set]], axis=0),
        ['confusion matrix', 'confusion matrix', 'accuracy']*1,
        [1, 2, '']*1
    ]
    idx = [
        np.concatenate([(i)*2) for i in classifier_names], axis=0),
        [1, 2] * len(classifier_names)
    ]
    results = pd.DataFrame(columns=cols, index = idx)
    results.index.names = ['classifier', 'actual']

    X = data_set.iloc[:, :-1].values
    y = data_set.iloc[:, -1].values

    for clf in classifier_:
        clf_name = str(clf.__class__).split('.')[1][-1][:-2]
        y_predict = cross_val_predict(clf, X, y)
        cm = confusion_matrix(y, y_predict)
        acc = accuracy_score(y, y_predict)
        results.loc[clf_name, (data_set.name, 'confusion matrix')] = cm
        results.loc[clf_name, (data_set.name, 'accuracy')] = acc

    return results

get_evaluation(exp2a, [des_tree, knn, gau_nav_bay, log_reg, svm_clf])
```

		exp2a		accuracy
		confusion matrix		
		1	2	
classifier	actual			
	DecisionTreeClassifier	1	550	0
	2	0	551	1.0
KNeighborsClassifier	1	550	0	1.0
	2	0	551	1.0
GaussianNB	1	0	550	0.02634
	2	522	29	0.02634
LogisticRegression	1	0	550	0.02634
	2	522	29	0.02634
SVC	1	220	330	0.499546

We can see, that Decision Tree and KNN are quite successfull. It has to be assumed that both classifiers are heavily overfitted. However, as the dataset is noise-free this is no real problem for this use case. GausianNB as well as Logistic Regression on the other hand perform incredibly bad. This is probably due to the fact that both classes themselves consist of two independent blocks each which are contrary to one another. However, the SVC employing a linear kernel is also able to separate the two classes in an acceptable fashion. Therefore it's somewhat surprising, that GausianNB and LogisticRegression perform this poorly. In Fact, the results for GausianNB and LogisticRegression are so bad that the replacement of a predicted class with the other class provides an almost perfect classifier. Given that observation it seems likely that there is some bug in the code. Sadly I've not been able to spot it yet.

## 2.2 (b)

Gaussian Naive Bayes, Logistic Regression and SVC show an accuracy that could be improved. I'll use GridSearchCV to search for better hyper-parameter.

### 2.2.1 Gaussian Naive Bayes

```
[ ]: X = exp2a.iloc[:, :-1].values
y = exp2a.iloc[:, -1].values

list_improved_clf = []
gs_gnb = GridSearchCV(gau_nav_bay, {'var_smoothing': np.logspace(300, 0, num=200)}, cv=10, n_jobs=-1)
gs_gnb.fit(X, y)
improved_clfs = [gs_gnb]
gs_gnb.best_params_
[ ]: {'var_smoothing': 1e+300}
```

### 2.2.2 Logistic Regression

```
[ ]: space = dict()
space['solver'] = ['newton-cg', 'lbfgs', 'liblinear']
space['penalty'] = ['none', 'l1', 'l2', 'elasticnet']
space['C'] = np.logspace(-5, 5, num=20)
model = LogisticRegression()

# disable warnings to increase readability. Some combinations of penalty and C values produce errors.
warnings.simplefilter(action='ignore', category=UserWarning)
warnings.simplefilter(action='ignore', category=sklearn.exceptions.FitFailedWarning)
```

```

gs_lr = RandomizedSearchCV(model, space, cv=10, n_jobs=-1, n_iter=500, u
↪scoring='accuracy', error_score=0)
gs_lr.fit(X,y)
improved_clfs.append(gs_lr)

warnings.simplefilter(action='default', category=UserWarning)
warnings.simplefilter(action='default', category=sklearn.exceptions.
↪FitFailedWarning)

gs_lr.best_params_

```

[ ]: {'solver': 'liblinear', 'penalty': 'l1', 'C': 1e-05}

### 2.2.3 SVC

```

[ ]: space = dict()
space['kernel'] = ['linear', 'poly', 'rbf', 'sigmoid']
scaler = StandardScaler()
X_transform = scaler.fit_transform(X)
gs_svm = GridSearchCV(svm_clf, param_grid=space, cv=10, scoring='accuracy', u
↪n_jobs=-1)
gs_svm.fit(X_transform,y)
improved_clfs.append(gs_svm)
gs_svm.best_params_

```

[ ]: {'kernel': 'poly'}

### 2.2.4 Accuracy comparison

```

[ ]: old_acc = (
    get_evaluation(exp2a, [gau_nav_bay, log_reg, svm_clf])
        .select_dtypes(object)
        .astype(float)
        .xs('accuracy', 1, 1)
        .groupby('classifier', 0)
        .agg('mean')
)

new_acc = (
    get_evaluation(exp2a, [i.best_estimator_ for i in improved_clfs])
        .select_dtypes(object)
        .astype(float)
        .xs('accuracy', 1, 1)
        .groupby('classifier', 0)
        .agg('mean')
).merge(old_acc, left_index=True, right_index=True)

```

```

new_acc.columns = ['new acc', 'old acc']

new_acc.assign(difference = lambda x: x['new acc'] - x['old acc'])

```

```
[ ]:          new acc   old acc   difference
classifier
GaussianNB      0.499546  0.026340  0.473206
LogisticRegression 0.499546  0.026340  0.473206
SVC            1.000000  0.499546  0.500454

```

it can be seen that the change of the hyperparameters has a tremendous effect on the accuracy of the classifiers.

### 2.3 (c) and (d)

```

[ ]: cols = [
    np.concatenate(([([i]*3) for i in ['exp2a', 'exp2b']], axis=0),
    ['confusion matrix', 'confusion matrix', 'accuracy']*2,
    ['positive', 'negative', '']*2
]
idx = [
    np.concatenate(([([str(i.__class__).split('.')[1][-1][:-2]]*2) for i in [
        des_tree, knn, gau_nav_bay, log_reg, svm_clf]], axis=0),
    ['positive', 'negative'] * 5
]
results_2_c_d = pd.DataFrame(columns=cols, index = idx)
results_2_c_d.index.names = ['classifier', 'actual']

for exp in [exp2a, exp2b]:
    X = exp.iloc[:, :-1].values
    y = exp.iloc[:, -1].values

    scaler = StandardScaler()
    X = scaler.fit_transform(X)

    for clf in [des_tree, knn, gau_nav_bay, log_reg, svm_clf]:
        name = str(clf.__class__).split('.')[1][-1][:-2]
        y_predict = cross_val_predict(clf, X, y)
        cm = confusion_matrix(y, y_predict)
        acc = accuracy_score(y, y_predict)
        results_2_c_d.loc[name, (exp.name, 'confusion matrix')] = cm
        results_2_c_d.loc[name, (exp.name, 'accuracy')] = acc
results_2_c_d

```

```
[ ]:           exp2a
confusion matrix
                           \
accuracy
```

		actual	positive	negative	
classifier	DecisionTreeClassifier	positive	550	0	1.0
		negative	0	551	1.0
KNeighborsClassifier		positive	550	0	1.0
		negative	0	551	1.0
GaussianNB		positive	0	550	0.02634
		negative	522	29	0.02634
LogisticRegression		positive	0	550	0.02634
		negative	522	29	0.02634
SVC		positive	220	330	0.499546
		negative	221	330	0.499546

		confusion matrix	accuracy	
		positive	negative	
classifier	actual			
DecisionTreeClassifier	positive	499	51	0.912807
	negative	45	506	0.912807
KNeighborsClassifier	positive	361	189	0.608538
	negative	242	309	0.608538
GaussianNB	positive	132	418	0.26158
	negative	395	156	0.26158
LogisticRegression	positive	175	375	0.331517
	negative	361	190	0.331517
SVC	positive	269	281	0.499546
	negative	270	281	0.499546

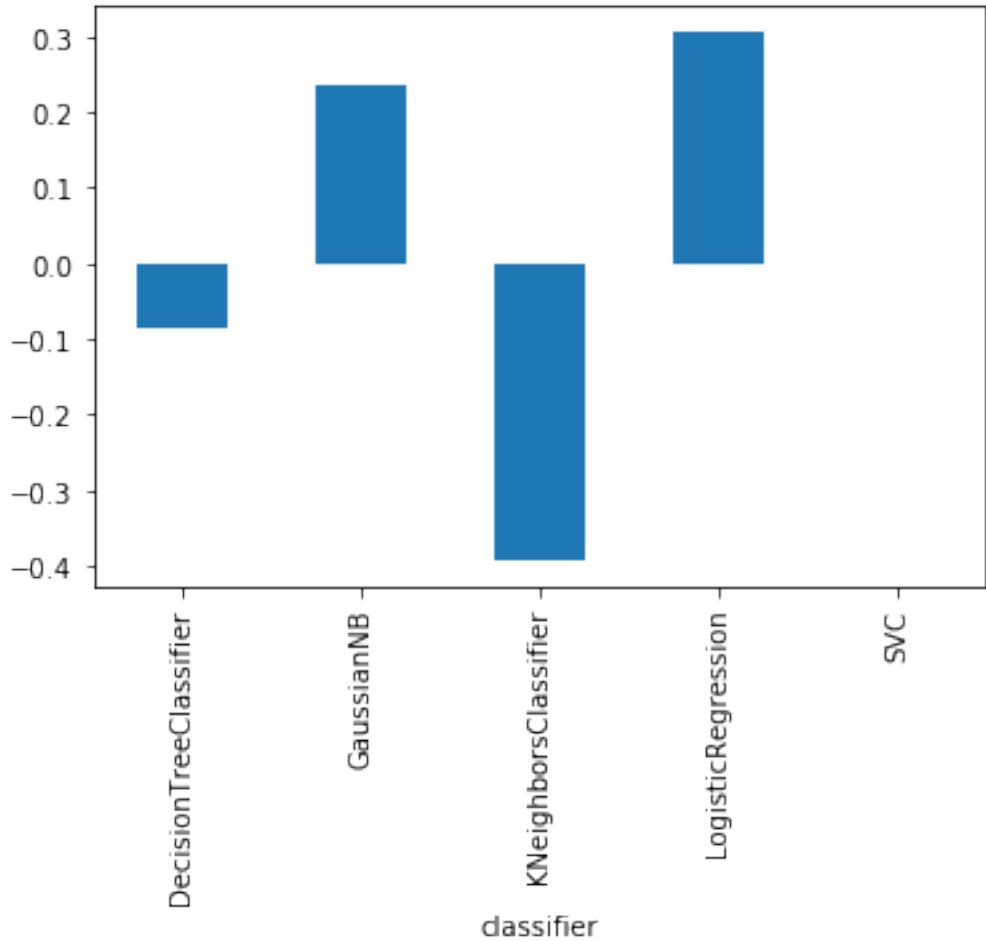
```
[ ]: accuracy_comp = (
    results_2_c_d
        .select_dtypes(object)
        .astype(float)
        .xs('accuracy', axis=1, level=1)
        .groupby('classifier', level=0, axis=0)
        .agg('mean')
        .assign(difference = lambda x: x['exp2b'] - x['exp2a'])
)
accuracy_comp['difference'].plot.bar()
accuracy_comp
```

```
[ ]: exp2a      exp2b difference

classifier
DecisionTreeClassifier 1.000000  0.912807 -0.087193
GaussianNB              0.026340  0.261580  0.235241
KNeighborsClassifier    1.000000  0.608538 -0.391462
LogisticRegression      0.026340  0.331517  0.305177
```

SVC

0.499546 0.499546 0.000000



Given the fact that the dataset was enriched by adding noisy features, sampled from a totally unrelated distribution it's surprising to see that in fact some classifiers worked even better than expected.

- Decision Trees
  - decrease can most likely be explained by noisy variables making it harder to generate linear decision boundaries parallel to axes. On the other hand that decrease is surprising, as the decision tree employs embedded feature selection, which should root out (at least part of) the noisy variables added to the set.
- Gaussian Naive Bayes
  - In 1.2(b) it was shown, that the performance of the Naive Bayes Classifier is influenced by the parameter ‘var\_smoothing’. Given the increased number of features and the non-linearity of the resulting data set it's surprising, that the NaiveBayes Classifier reaches a better result for exp2b, than for exp2a
- Logistic Regression
  - In 1.2(b) it was shown, that the Logistic Regression can be heavily optimized using

many different parameters. The implementation as shown in Table 1 is simply not optimal. However, the increase in performance is still surprising, given that the added noisy features dont contribute to the linearity of the data-set it's surprising that the logistic regression yields better results for exp2b than for exp2a.

- Nearest Neighbor
  - the increased number of features invokes the ‘cures of dimensionality’. This means that the density of features is decreased, which leads to worse generalization performance.
- Support Vector Machines
  - 1.2(b) has shown that SVM is highly dependent on the kernel function. As this has not changed it's n suprising, that the results are also unchangeotd.

### 3 Assignment 1.3

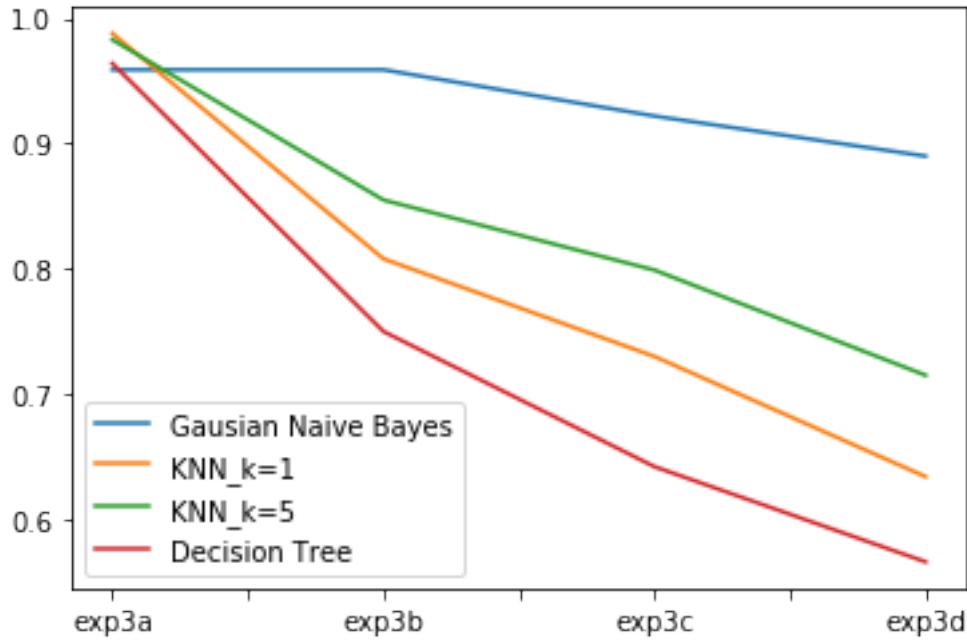
```
[ ]: exp3a = pd.read_csv('exp3a.csv')
exp3a.name = 'exp3a'
exp3b = pd.read_csv('exp3b.csv')
exp3b.name = 'exp3b'
exp3c = pd.read_csv('exp3c.csv')
exp3c.name = 'exp3c'
exp3d = pd.read_csv('exp3d.csv')
exp3d.name = 'exp3d'
```

```
[ ]: des_tree = DecisionTreeClassifier(criterion='entropy')
knn = KNeighborsClassifier()
gau_nav_bay = GaussianNB()
log_reg = LogisticRegression()
svm_clf = SVC(kernel='linear')
```

#### 3.1 (a)

```
[ ]: idx = ['Gausian Naive Bayes', 'KNN_k=1', 'KNN_k=5', 'Decision Tree']
cols = ['exp3a', 'exp3b', 'exp3c', 'exp3d']
results_1_3_a = pd.DataFrame(index=idx, columns=cols)
for name, clf in zip(['Gausian Naive Bayes', 'KNN_k=1', 'KNN_k=5', 'Decision Tree'],
                     [GaussianNB(), KNeighborsClassifier(n_neighbors=1),
                      KNeighborsClassifier(n_neighbors=5),
                      DecisionTreeClassifier(criterion='entropy')]):
    for exp in [exp3a, exp3b, exp3c, exp3d]:
        X = exp.iloc[:, :-1]
        y = exp.iloc[:, -1]
        y_predict = cross_val_predict(clf, X, y, cv=10, n_jobs=-1)
        acc = accuracy_score(y, y_predict)
        results_1_3_a.loc[name, exp.name] = acc
results_1_3_a.T.plot.line()
results_1_3_a
```

```
[ ]:          exp3a  exp3b  exp3c  exp3d
Gausian Naive Bayes  0.959  0.959  0.922  0.89
KNN_k=1             0.988  0.808  0.73   0.634
KNN_k=5             0.983  0.855  0.799  0.715
Decision Tree        0.964  0.75   0.642  0.566
```

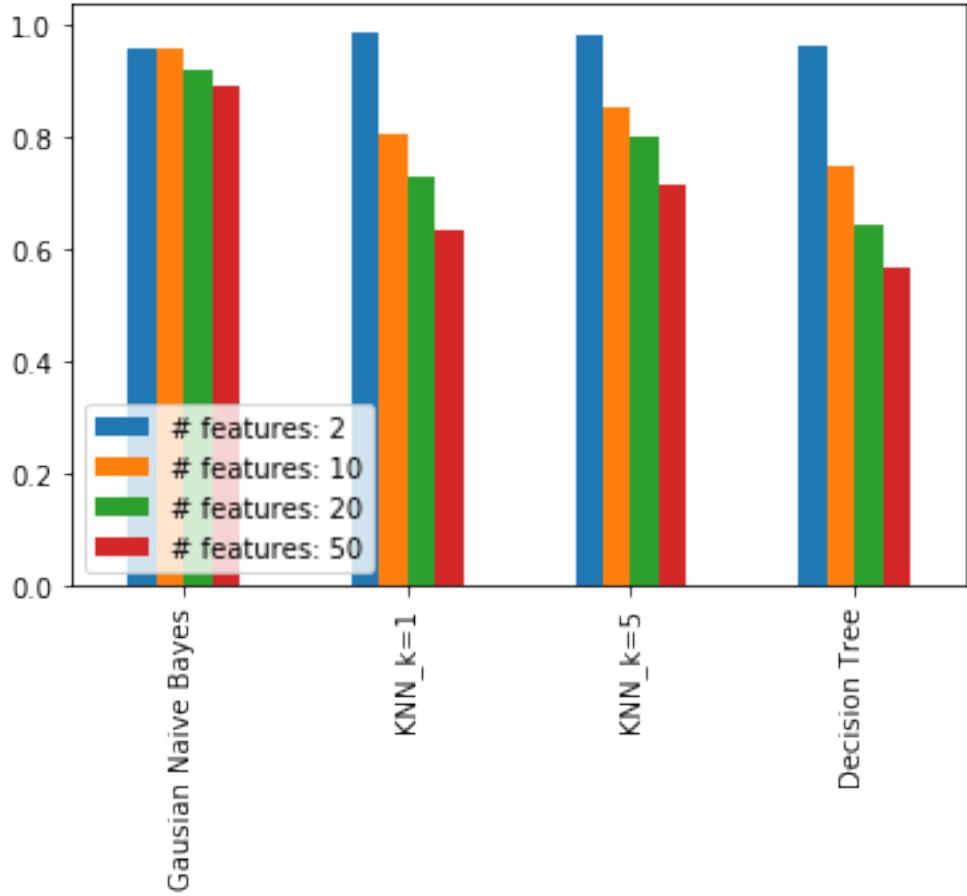


We can see that all classifiers have the highest accuracy for exp3a and the lowest for exp3d.

### 3.2 (b)

```
[ ]: results_1_3_a.plot.bar()
plt.legend(labels = [f'# features: {exp.iloc[:, :-1].shape[1]}' for exp in
                     [exp3a, exp3b, exp3c, exp3d]], loc=3)
results_1_3_a.columns = [f'# features: {exp.iloc[:, :-1].shape[1]}' for exp in
                        [exp3a, exp3b, exp3c, exp3d]]
results_1_3_a
```

```
[ ]:          # features: 2 # features: 10 # features: 20 # features: 50
Gausian Naive Bayes      0.959      0.959      0.922      0.89
KNN_k=1                  0.988      0.808      0.73       0.634
KNN_k=5                  0.983      0.855      0.799      0.715
Decision Tree            0.964      0.75       0.642      0.566
```



To evaluate this correctly we have to take into account that the number of features increases from exp3a to exp3d. As we can clearly see the increase of features leads to a decrease in accuracy. This effect is of different strength for different classifiers. The best performing classifier is Gausian Naive Bayes, which only decreases from 0.959 for 2 features to 0.89 for 50 features. k-NearestNeighbors on the other hand perform worse for increasing numbers of features. This can be very well explained with the curse of dimensionality. This describes the fact that data is less dense the more dimensions are taken into account. It's noteworthy however, that a smaller k ( $k=1$ ) performs better for low dimensions than a higher k ( $k=5$ ). For higher dimensions however this effect is reversed. So KNN\_k=1 declines from 0.988 to 0.634, whereas the KNN\_k=5 starts lower at 0.983, but ends higher at 0.715. The decision tree sees the highest drop in accuracy (from 0.966 to 0.554). This can be explained with the fact, that the data is distributed in a way, that's very hard to capture for a decision tree. For higher dimensions this effect increases, meaning that the individual decision trees are heavily overfitted, so they don't generalize very well.