

# Lab3

October 6, 2021

```
[ ]: import pandas as pd
import numpy as np
from sklearn.ensemble import RandomForestClassifier, AdaBoostClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import SVC
from sklearn.model_selection import cross_val_predict, cross_val_score, GridSearchCV, RandomizedSearchCV, StratifiedKFold, StratifiedGroupKFold
from sklearn.metrics import confusion_matrix, accuracy_score, precision_recall_fscore_support, recall_score, roc_auc_score, classification_report
from sklearn.preprocessing import StandardScaler, MinMaxScaler
import matplotlib.pyplot as plt
import scipy
from sklearn.pipeline import make_pipeline, Pipeline
from sklearn.feature_selection import chi2, SelectKBest, RFECV, SelectFromModel, SequentialFeatureSelector
from sklearn.inspection import PartialDependenceDisplay, partial_dependence
from sklearn.tree import DecisionTreeClassifier
from sklearn.decomposition import PCA
from pdpbox import pdp, get_dataset, info_plots
```

```
[ ]: df = pd.read_csv('caravan.csv')
```

One-hot-encode categorical features - which features are categorical was derived from the feature description in the assignment

```
[ ]: cat_cols = ['Customer Subtype', 'Customer main type', 'Roman catholic', 'Contribution private third party insurance']
df[cat_cols] = df[cat_cols].astype(object)
df = pd.get_dummies(df)
```

```
[ ]: y = df['CARAVAN POLICY']
X = df.drop('CARAVAN POLICY', axis=1)
```

```
[ ]: df['CARAVAN POLICY'].value_counts()
```

```
[ ]: 0      5474
     1      348
```

```
Name: CARAVAN POLICY, dtype: int64
```

As can be seen the whole dataset is very imbalanced. Therefore accuracy is not that good a metric to decide the goodness of the classifier. Instead I'll use the f1\_weighted as the scoring metric.

```
[ ]: seq_feat = SequentialFeatureSelector(estimator=DecisionTreeClassifier(),  
    ↪n_features_to_select=10, cv=None, n_jobs=-1)  
seq_feat.fit(X, y)  
X_new = seq_feat.fit_transform(X,y)  
  
X_new = pd.DataFrame(X_new, columns=seq_feat.get_feature_names_out())
```


```
[ ]: param_grid = dict(C = np.logspace(-4,4,10))  
clf = GridSearchCV(SVC(class_weight='balanced', probability=True),  
    ↪param_grid=param_grid, n_jobs=-1, scoring='f1_weighted')  
clf.fit(X_new, y)  
clf = clf.best_estimator_  
y_pred = clf.predict(X_new)  
print(classification_report(y, y_pred))
```

	precision	recall	f1-score	support
0	0.95	0.60	0.74	5474
1	0.08	0.55	0.14	348
accuracy			0.60	5822
macro avg	0.52	0.58	0.44	5822
weighted avg	0.90	0.60	0.70	5822

```
[ ]: explain_df = X_new.merge(y, left_index=True, right_index=True)  
  
for i in range(10):  
    pdp_goals = pdp.pdp_isolate(model=clf, dataset=explain_df, model_features=  
        ↪explain_df.columns[:-1], feature=explain_df.columns[i])  
  
    pdp.pdp_plot(pdp_goals, explain_df.columns[i])  
    plt.show()
```


PDP for feature "Number of houses"

Number of unique grid points: 2




PDP for feature "Avg size household"

Number of unique grid points: 5




PDP for feature "Contribution third party insurance (firms)"

Number of unique grid points: 2




PDP for feature "Contribution motorcycle/scooter policies"

Number of unique grid points: 2




PDP for feature "Contribution lorry policies"

Number of unique grid points: 2




PDP for feature "Contribution agricultural machines policies"

Number of unique grid points: 2




PDP for feature "Contribution moped policies"

Number of unique grid points: 2




PDP for feature "Contribution private accident insurance policies"

Number of unique grid points: 2




PDP for feature "Contribution bicycle policies"

Number of unique grid points: 2



PDP for feature "Number of social security insurance policies"

Number of unique grid points: 2



These partial dependence plots explain the 10 selected features and their influence on the likelihood that a customer will end up buying a caravan insurance.

# 1 Assignment 2

```
[ ]: df_test = pd.read_csv('caravanTest.csv')
```

```
[ ]: X_test = df_test[X_new.columns]
y_test = df_test['CARAVAN POLICY']
y_pred_test = clf.predict(X_test)
print(classification_report(y_test, y_pred_test))
```

	precision	recall	f1-score	support
0	0.94	0.57	0.71	3762
1	0.06	0.43	0.10	238
accuracy			0.56	4000
macro avg	0.50	0.50	0.41	4000
weighted avg	0.89	0.56	0.68	4000

```
[ ]: results = pd.DataFrame(clf.predict_proba(X_test), columns = ['prediction_0',  
→'prediction_1'])
```

```
[ ]: results['prediction_1'].nlargest(800).index
```

```
[ ]: Int64Index([1726, 3931, 865, 1069, 2128, 3897, 68, 229, 1393, 1760,  
..  
1749, 1751, 1752, 1754, 1755, 1759, 1764, 1765, 1766, 1767],  
dtype='int64', length=800)
```

This returns the indices of the 800 candidates with the highest predicted probability of buying a caravan policy.